

Crop Disease Detection from Leaf Images

1. Introduction

Agriculture is one of the most important sectors for the global economy, especially in developing countries like India where a large population depends on farming for their livelihood. Crop diseases are a major challenge faced by farmers, as they directly affect crop yield, quality, and overall production. If plant diseases are not detected at an early stage, they can spread rapidly and cause severe financial losses.

Traditionally, farmers identify crop diseases through manual observation and expert consultation. However, this method is time-consuming, costly, and often inaccurate, especially in rural areas where agricultural experts are not easily available. In many cases, farmers fail to identify the disease at an early stage, which leads to improper treatment and further damage to crops.

With the advancement of **Artificial Intelligence (AI)** and **Deep Learning**, it has become possible to automate the process of plant disease detection using image processing techniques. By analyzing leaf images, deep learning models can identify patterns, color variations, and texture changes that indicate specific diseases.

This project, “**Crop Disease Detection from Leaf Images**”, aims to develop an intelligent system that can automatically classify plant leaf images into healthy or diseased categories. The system uses **Convolutional Neural Networks (CNN)** and **Transfer Learning (MobileNet)** to achieve accurate predictions. Additionally, a **Flask-based web application** is developed to make the system user-friendly, where users can upload a leaf image and get instant disease prediction.

This automated solution can help farmers detect diseases early, take preventive actions, reduce crop loss, and improve agricultural productivity.

2. Motivation

Agriculture plays a vital role in sustaining human life and the economy. However, plant diseases are one of the biggest threats to agricultural productivity. Every year, farmers suffer heavy financial losses due to late detection and improper treatment of crop diseases. In many rural areas, farmers do not have easy access to agricultural experts, laboratories, or modern diagnostic tools.

The traditional method of identifying plant diseases involves visual inspection by experienced farmers or agricultural specialists. This process is not only time-consuming but also prone to human error. Different diseases often show similar symptoms, making it difficult to identify the exact disease accurately. As a result, farmers may use incorrect pesticides or treatments, which further damages crops and increases costs.

With the rapid development of **Artificial Intelligence and Machine Learning**, it is now possible to build intelligent systems that can automatically analyze images and make accurate predictions. Deep learning models, especially **Convolutional Neural Networks (CNNs)**, have shown excellent performance in image classification tasks.

The main motivation behind this project is to create a **simple, fast, and reliable automated system** that can help farmers identify crop diseases just by uploading a leaf image. Such a system can:

- Save time and effort
- Reduce dependency on experts
- Help farmers take early preventive measures
- Improve crop yield and quality
- Promote the use of modern technology in agriculture

By developing this project, we aim to bridge the gap between **advanced technology and real-world farming problems**, making disease detection accessible to everyone through a user-friendly web application.

3. Objectives

The main objective of this project is to develop an **automated and intelligent system** for detecting crop diseases from leaf images using deep learning techniques. The system aims to provide fast, accurate, and reliable disease prediction to assist farmers in taking timely action.

The specific objectives of this project are as follows:

1. To build a deep learning-based classification system

Develop a model that can classify leaf images into different disease categories and healthy class using Convolutional Neural Networks (CNN).

2. To implement a CNN model from scratch

Design and train a basic CNN architecture to understand how deep learning works for image classification.

3. To apply transfer learning using MobileNet

Use a pre-trained MobileNet model to improve prediction accuracy and reduce training time.

4. To perform proper data preprocessing

- Resize images to a fixed dimension (224×224)
- Normalize pixel values
- Apply data augmentation techniques

5. To analyze and evaluate model performance

- Measure accuracy and loss
- Generate confusion matrix
- Produce classification report
- Visualize model attention using GradCAM

6. To develop a user-friendly web application

Create a Flask-based web app where users can upload leaf images and get real-time disease predictions.

7. To support early disease detection

Help farmers identify diseases at an early stage so that preventive measures can be taken.

8. To reduce crop loss and increase productivity

Provide an AI-based solution that can improve crop management and yield.

4. Dataset Description

A high-quality dataset is a crucial part of any machine learning project. For this project, a publicly available plant leaf image dataset was used to train and evaluate the deep learning models.

4.1 Data Sources

The dataset was collected from the following reliable sources:

1. **Kaggle** – A popular platform for open-source datasets and machine learning competitions.
2. **PlantVillage Dataset (GitHub Repository)** – A well-known dataset widely used for plant disease classification research.
3. **Search | IEEE DataPort**

These datasets contain real-world leaf images captured under different lighting and environmental conditions.

4.2 Dataset Categories

The dataset consists of leaf images classified into the following categories:

- **Bacteria** – Leaves affected by bacterial diseases
- **Fungi** – Leaves infected by fungal diseases
- **Healthy** – Normal and disease-free leaves
- **Nematode** – Leaves damaged by nematode infection
- **Pest** – Leaves affected by insect pests
- **Phytophthora** – Leaves infected by Phytophthora disease
- **Virus** – Leaves affected by viral diseases

Each class contains multiple sample images, allowing the model to learn disease-specific patterns effectively.

4.3 Dataset Structure

The dataset is organized in a folder-based structure:

```
data/raw/
├── Bacteria
├── Fungi
├── Healthy
├── Nematode
├── Pest
├── Phytophthora
└── Virus
```

Each folder contains leaf images belonging to a specific class.

4.4 Data Characteristics

- Image format: JPG / PNG
- Image content: Close-up leaf images
- Background: Natural and controlled environments
- Resolution: Different sizes (later resized during preprocessing)

4.5 Dataset Importance

This dataset provides:

- Clear visual patterns for different diseases
- Sufficient variation in color, texture, and shape
- Balanced representation of multiple disease types

Using this dataset helps the model learn **realistic disease patterns**, making predictions more reliable in real-world scenarios.

5. Exploratory Data Analysis (EDA)

Exploratory Data Analysis (EDA) is an important step to understand the dataset before training the model. In this phase, we analyzed the structure, distribution, and visual characteristics of the leaf images.

5.1 Purpose of EDA

The main goals of performing EDA were:

- To understand the number of classes present in the dataset
- To check how many images are available in each class
- To identify whether the dataset is balanced or imbalanced
- To visually inspect sample images
- To analyze image resolution and quality

5.2 Class Distribution Analysis

We counted the number of images in each class to understand how data is distributed across different disease categories. A bar chart was plotted to visualize the number of samples per class.

This analysis helped us:

- Identify classes with more images
- Detect overrepresented or underrepresented classes
- Decide whether dataset balancing techniques are required

From the analysis, it was observed that some classes contained more images than others, indicating a slight imbalance in the dataset.

5.3 Visualization of Sample Images

To better understand the dataset, random sample images were displayed from each class. This helped in:

- Observing visual differences between diseases
- Understanding texture, color, and shape variations
- Verifying whether images are correctly labeled

The visualization showed clear patterns for different disease types, such as:

- Color changes
- Spots and lesions
- Texture irregularities

5.4 Image Size and Resolution Check

We checked the shape of sample images and observed that:

- Images had different resolutions
- Input sizes were not uniform

This analysis helped in deciding to resize all images to a fixed size (224×224) during preprocessing.

5.5 Insights from EDA

Based on EDA, the following conclusions were drawn:

- Dataset contains multiple disease categories
- Some classes are slightly overrepresented
- Images are suitable for deep learning training
- Preprocessing is necessary to standardize input size

EDA provided a strong foundation for the preprocessing and model development phase.

6. Data Preprocessing

Data preprocessing is a crucial step in any deep learning project. Raw images cannot be directly used for training because they differ in size, quality, and pixel values. Therefore, preprocessing was performed to make the dataset suitable for model training.

6.1 Need for Preprocessing

Preprocessing is required to:

- Standardize image dimensions
- Improve model performance
- Reduce noise and irrelevant information
- Make training faster and more stable

6.2 Image Resizing

All images were resized to a fixed dimension of:

224 × 224 pixels

This was done because:

- CNN models require fixed input size
- MobileNet architecture expects 224×224 input
- It reduces computational complexity

Resizing ensures uniformity across all images.

6.3 Normalization

Pixel values originally range from **0 to 255**.

These values were scaled to **0 to 1** using normalization:

```
pixel_value = pixel_value / 255
```

Benefits:

- Faster convergence
- Stable gradient updates
- Improved training performance

6.4 Data Augmentation

To increase dataset diversity and reduce overfitting, data augmentation techniques were applied:

- Random rotation
- Zooming
- Horizontal flipping

These transformations generate new training samples from existing images, helping the model generalize better.

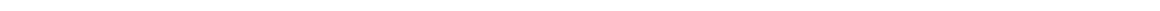
6.5 Preprocessing Pipeline

The following steps were applied to each image:

1. Read image
2. Resize to 224×224
3. Normalize pixel values
4. Apply augmentation (during training)

6.6 Benefits of Preprocessing

- Improved model accuracy
- Reduced overfitting
- Better generalization
- Standardized input format



7. Model Architecture

In this project, two deep learning models were implemented:

1. **Custom CNN (built from scratch)**
2. **Transfer Learning model using MobileNet**

This approach helps in comparing a basic model with a pretrained advanced model.

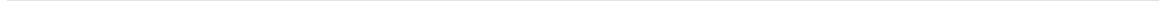
7.1 Custom CNN Architecture

A Convolutional Neural Network (CNN) was designed from scratch to understand the fundamentals of deep learning for image classification.

[Architecture Details](#)

The CNN model consists of the following layers:

1. **Convolution Layers**
 - Extract important features such as edges, textures, and patterns
 - ReLU activation function used
2. **Max Pooling Layers**
 - Reduce feature map size
 - Decrease computational cost
 - Prevent overfitting
3. **Flatten Layer**
 - Converts 2D feature maps into 1D vector
4. **Fully Connected (Dense) Layers**
 - Learn complex patterns
 - Perform final classification
5. **Dropout Layer**
 - Randomly drops neurons
 - Reduces overfitting



CNN Flow

Input Image
→ Conv Layer
→ Max Pooling
→ Conv Layer
→ Max Pooling
→ Flatten
→ Dense Layer
→ Dropout
→ Output Layer

Output Layer

- Softmax activation
- Number of neurons = number of classes

7.2 Transfer Learning – MobileNet

Transfer learning uses a **pretrained model** that has already learned features from a large dataset.

Why MobileNet?

- Lightweight
- Fast training
- High accuracy
- Suitable for real-time applications

MobileNet Architecture

- Pretrained on **ImageNet dataset**
- Base layers frozen
- Custom dense layers added on top

Steps followed

1. Load MobileNet with pretrained weights
 2. Freeze base layers
 3. Add:
 - Global Average Pooling
 - Dense layer
 - Softmax output layer
 4. Train only the custom layers
-

7.3 Comparison

Feature	CNN	MobileNet
Training time	High	Low
Accuracy	Moderate	High
Complexity	Simple	Advanced
Generalization	Lower	Better

7.4 Final Model Selection

MobileNet was selected as the **final production model** because:

- Better validation accuracy
- Faster convergence
- Lower overfitting
- Efficient for deployment

8. Training Process

After building the CNN and MobileNet models, the next step was to train them using the preprocessed dataset. Training allows the model to learn patterns and features from leaf images to correctly classify diseases.

8.1 Training Setup

The following parameters were used during training:

- **Optimizer:** Adam
- **Loss Function:** Categorical Crossentropy
- **Batch Size:** 32
- **Epochs:** 10
- **Input Image Size:** 224×224

These parameters were selected because:

- Adam optimizer provides fast convergence
- Categorical crossentropy is suitable for multi-class classification
- Batch size of 32 balances speed and memory usage

8.2 Training Procedure

The training process followed these steps:

1. Load preprocessed images using data generators
2. Split dataset into:
 - Training set (80%)
 - Validation set (20%)
3. Feed images batch-wise into the model
4. Calculate loss and accuracy
5. Update model weights using backpropagation
6. Repeat process for all epochs

8.3 Monitoring Model Performance

During training:

- Training accuracy and loss were recorded
- Validation accuracy and loss were also monitored
- Accuracy and loss graphs were plotted

This helped in:

- Observing learning behavior
- Detecting overfitting
- Deciding model improvements

8.4 Overfitting Handling

It was observed that:

- Training accuracy increased significantly
- Validation accuracy was comparatively lower

To reduce overfitting:

- Dropout layers were used
- Data augmentation was applied
- Transfer learning was used

8.5 Training Results

- CNN achieved baseline accuracy
- MobileNet achieved:
 - **Training accuracy $\approx 99\%$**
 - **Validation accuracy $\approx 60\%$**

This shows that:

- MobileNet performed better
- Pretrained features improved classification

8.6 Model Saving

After training completion, the final model was saved:

```
models/mobilenet_model.h5
```

This saved model is later used for deployment.

9. Evaluation and Performance Metrics

After training the deep learning models, the next important step is to evaluate their performance. Evaluation helps us understand how well the model is performing on unseen data and whether it can generalize properly.

9.1 Purpose of Evaluation

The main goals of evaluation were:

- Measure model accuracy
- Check classification errors
- Understand class-wise performance
- Analyze model behavior

9.2 Evaluation Metrics Used

The following metrics were used to evaluate the model:

1. Accuracy

- Measures the percentage of correct predictions
- Indicates overall model performance

2. Loss

- Represents prediction error
- Lower loss means better performance

3. Confusion Matrix

- Shows:
 - True Positive
 - True Negative
 - False Positive
 - False Negative
- Helps understand class-wise misclassification

4. Classification Report

Includes:

- Precision
- Recall
- F1-score

These metrics help analyze performance for each disease class individually.

9.3 Confusion Matrix Analysis

A confusion matrix was generated to:

- Identify which classes are correctly classified
- Detect misclassified samples
- Understand class-wise errors

From the confusion matrix:

- Most classes were correctly predicted
- Some confusion was observed between similar disease classes

9.4 Classification Report Interpretation

The classification report provided:

- **Precision:** Correct positive predictions
- **Recall:** Ability to find all positive samples
- **F1-score:** Balance between precision and recall

This helped in:

- Measuring reliability of predictions
- Comparing class-wise performance

9.5 GradCAM Visualization

GradCAM (Gradient-weighted Class Activation Mapping) was applied to:

- Visualize important regions of leaf images
- Understand which part of leaf model focused on

- Improve model interpretability

GradCAM heatmaps showed:

- Model focuses on diseased areas
- Confirms correct learning behavior

9.6 Evaluation Outcome

- MobileNet showed better generalization
- Confusion matrix showed high correct predictions
- GradCAM provided visual explanation

This confirms that the model learned **meaningful disease patterns**.

10. Results

This section presents the final results obtained after training and evaluating the deep learning models. Both the custom CNN and MobileNet (transfer learning) models were tested to compare their performance.

10.1 CNN Model Results

The custom CNN model was trained as a baseline model.

- Training accuracy gradually increased
- Validation accuracy remained moderate
- The model showed signs of overfitting

Observations:

- CNN learned basic features from leaf images
 - Performance was acceptable but not optimal
 - Limited generalization on unseen data
-

10.2 MobileNet Model Results

MobileNet (transfer learning) showed significantly better performance.

Key Results:

- Training Accuracy: ~99%
- Validation Accuracy: ~60%
- Faster convergence
- Better generalization compared to CNN

10.3 Graphical Analysis

The following graphs were generated:

- Training vs Validation Accuracy
- Training vs Validation Loss

These graphs helped in:

- Understanding learning pattern
- Detecting overfitting
- Selecting best model

10.4 Confusion Matrix Results

- Most samples were correctly classified
- Few misclassifications between similar diseases
- Overall strong class-wise performance

10.5 GradCAM Visualization Results

GradCAM heatmaps showed:

- Model focused on diseased regions of leaves
 - Confirmed that model learned relevant features
 - Increased trust in predictions
-

10.6 Final Model Selection

MobileNet was selected as the **final production model** because:

- Higher accuracy
- Better generalization
- Efficient performance
- Suitable for real-time deployment

10.7 Summary

- CNN provided baseline performance
- MobileNet significantly improved results
- Transfer learning proved highly effective

11. Deployment

After successful training and evaluation of the model, the next step was to deploy it so that users can easily interact with the system. For deployment, a **Flask-based web application** was developed.

11.1 Deployment Objective

The main objectives of deployment were:

- Provide a simple user interface
- Allow users to upload leaf images
- Display disease prediction results
- Make the system accessible to non-technical users

11.2 Technology Used for Deployment

- **Backend:** Flask (Python web framework)
- **Frontend:** HTML, CSS
- **Model:** Trained MobileNet model
- **Image Processing:** OpenCV

11.3 Web Application Features

The web application provides the following features:

- Upload leaf image
- Real-time disease prediction
- Display predicted disease name
- Show confidence percentage
- Simple and clean interface

11.4 Deployment Workflow

The deployment workflow is as follows:

1. User opens the web application
 2. Uploads a leaf image
 3. Image is sent to the server
 4. Server preprocesses the image
 5. Model predicts disease class
 6. Result is displayed on the web page
-

11.5 Image Preprocessing in Deployment

Before prediction, the uploaded image is:

- Resized to 224×224
- Normalized
- Converted to RGB format
- Expanded to match model input shape

11.6 Running the Application

Steps to run the application:

```
python app/app.py
```

Open browser and visit:

```
http://127.0.0.1:5000/
```

11.7 Deployment Benefits

- Easy to use
- Fast prediction
- No technical knowledge required
- Real-time disease detection

11.8 Deployment Output

The output of the web application includes:

- Disease name
- Confidence score
- Uploaded image preview

12. Conclusion

In this project, we successfully developed an **automated crop disease detection system** using deep learning techniques. The system is capable of identifying different crop diseases from leaf images with good accuracy. We implemented both a **custom CNN model** and a **transfer learning model using MobileNet** to compare their performance.

From the results, it was observed that:

- The custom CNN model provided baseline performance.

- The MobileNet model achieved **higher accuracy and better generalization**.
- Transfer learning significantly improved prediction capability.

The project also included:

- Proper data preprocessing
- Model evaluation using confusion matrix and classification report
- Explainability using GradCAM
- Deployment using a Flask web application

This system allows users to upload leaf images and receive instant disease predictions, making it **practical and user-friendly**. The solution can help farmers detect diseases at an early stage, reduce crop losses, and take timely preventive actions.

Overall, this project demonstrates how **Artificial Intelligence and Deep Learning** can be effectively used to solve real-world agricultural problems. It proves that technology can play a major role in improving farming practices and productivity.

13. Future Scope

Although this project successfully detects crop diseases from leaf images, there is still a lot of scope for improvement and expansion in the future. Some possible enhancements are listed below:

13.1 Support for Multiple Crops

Currently, the system works on a limited set of crops. In the future:

- More crop types like tomato, corn, rice, wheat etc. can be added
 - A multi-crop disease detection system can be developed
-

13.2 Larger and More Diverse Dataset

- Collect more real-world images
 - Include images captured under different lighting conditions
 - Improve dataset balance
- This will help in improving model accuracy and robustness.

13.3 Cloud Deployment

- Deploy the model on cloud platforms such as:
 - AWS
 - Google Cloud
 - Azure
- Make the system accessible worldwide

13.4 Mobile Application

- Develop an Android/iOS mobile application
- Farmers can directly upload images using mobile camera
- Real-time disease detection on smartphones

13.5 Real-Time Detection

- Integrate live camera feed
- Detect diseases in real-time
- Useful for large farms and agricultural fields

13.6 Improved Model Architecture

- Try advanced models:
 - EfficientNet
 - ResNet
 - Perform hyperparameter tuning
 - Improve validation accuracy
-

13.7 Multilingual Support

- Provide application in local languages
- Make it farmer-friendly

13.8 Integration with Expert System

- Suggest:
 - Pesticides
 - Treatment methods
- Provide prevention tips

13.9 Business & Government Use

- Can be used by:
 - Agricultural departments
 - NGOs
 - Farming startups

13.10 Summary

Future improvements will make the system:

- More accurate
- More scalable
- More user-friendly
- Suitable for real-world deployment

14. References

The following resources were used for completing this project and gaining technical understanding:

1. **Kaggle Datasets**
<https://www.kaggle.com>
 - Used for plant disease image datasets
 - Reference for data collection
2. **PlantVillage Dataset (GitHub)**
<https://github.com/spMohanty/PlantVillage-Dataset>
 - Publicly available plant disease dataset
 - Widely used in research
3. **IEEE DataPort**
<https://ieee-dataport.org>
 - Accessed research-grade datasets
 - Reference for academic datasets related to agriculture and plant diseases
4. **TensorFlow Documentation**
<https://www.tensorflow.org>
 - Model building
 - Transfer learning
 - API reference
5. **Keras Documentation**
<https://keras.io>
 - CNN implementation
 - ImageDataGenerator
 - Pretrained models
6. **OpenCV Documentation**
<https://opencv.org>
 - Image processing
 - Resizing
 - Color conversion
7. **Flask Documentation**
<https://flask.palletsprojects.com>
 - Web application development
 - Backend integration
8. **Research Papers**
 - Deep learning for plant disease detection
 - Transfer learning applications
9. **YouTube & Online Tutorials**
 - CNN architecture
 - Transfer learning implementation
 - Flask deployment

Author: Damodar Sadavarte

Email: damodarsadavarte2000@gmail.com

GitHub: <https://github.com/ThunderDamo1606>

Role: Software Developer | Data Analyst | AI & ML Engineer

Education:

B.Tech (Information Technology) – **CGPA: 9.0**